

Matrix Methods in Machine Learning

Lecture Notes

Rebekah Dix

October 5, 2018

Contents

1	Elements of Machine Learning	2
2	Linear Algebra Review	2
2.1	Products	2
2.2	Linear Independence	3
3	Linear Systems and Vector Norms	4
3.1	Vector Norms	4
4	Least Squares	5
4.1	Geometric Approach	6
4.2	Vector Calculus Approach	7
4.2.1	Review of Vector Calculus	7
4.2.2	Application to Least Squares	8
4.3	Positive Definite Matrices	8
4.4	Subspaces	9
4.5	Least Squares with Orthonormal Basis for Subspace	9
4.5.1	Orthogonal Matrices and Orthonormal Basis	9
4.5.2	Back to LS	10
4.5.3	Gram-Schmidt Orthogonalization Algorithm	10
5	Least Squares Classification	11
6	Tikhonov Regularization/Ridge Regression	12
6.1	Tikhonov Regularization Derivation	12
6.1.1	Derivation with Vector Calculus	12
6.1.2	Alternative Derivation	13
7	Singular Value Decomposition	14

1 Elements of Machine Learning

1. Collect data
2. Preprocessing: changing data to simplify subsequent operations without losing relevant information.
3. Feature extraction: reduce raw data by extracting features or properties relevant to the model.
4. Generate training samples: a large collection of examples we can use to learn the model.
5. Loss function: To learn the model, we choose a loss function (i.e. a measure of how well a model fits the data)
6. Learn the model: Search over a collection of candidate models or model parameters to find one that minimizes the loss on training data.
7. Characterize generalization error (the error of our predictions on new data that was not used for training).

2 Linear Algebra Review

2.1 Products

Inner products:

$$\langle x, w \rangle = \sum_{j=1}^p w_j x_j = x^T w = w^T x \quad (1)$$

Thus this inner product is a weighted sum of the elements of x .

Matrix-vector multiplication:

$$Xw = \begin{bmatrix} -x_1^T- \\ -x_2^T- \\ \vdots \\ -x_n^T- \end{bmatrix} w = \begin{bmatrix} x_1^T w \\ x_2^T w \\ \vdots \\ x_n^T w \end{bmatrix} \quad (2)$$

Matrix-matrix multiplication:

Example 1. Let $X \in \mathbb{R}^{n \times p}$, n movies, p people. $T \in \mathbb{R}^{n \times r}$, and $W \in \mathbb{R}^{r \times p}$. We can think of T as the taste profiles of r representative customers and W as the weights on each representative profile (there will be one set of weights for each customer). Suppose we have two representative taste profiles (i.e. an action lover and a romance lover). Then w will be a 2-vector containing the weights of on the two representative taste profiles. Then

Tw is the expected preferences of a customer who weights the representative taste profiles of T with the weights given in w .

Now we can think about the full matrix product $X = TW$

$$X = TW \implies X_{ij} = \langle \text{ith row of } T, j\text{th column of } W \rangle \quad (3)$$

- The j th column of X is a weighted sum of the columns of T , where the j th column of W tells us the weights.

$$x_j = Tw_j \quad (4)$$

That is, the tastes (preferences) of the j th customer.

- The i th row of X is $x_i^T = t_i^T W$ where t_i^T is the i th row of T . This gives us how much each customer likes movie i .

Inner product representation:

$$TW = \begin{bmatrix} -t_1^T- \\ -t_2^T- \\ \vdots \\ -t_n^T- \end{bmatrix} \begin{bmatrix} | & | & & | \\ w_1 & w_2 & \dots & w_p \\ | & | & & | \end{bmatrix} = \begin{bmatrix} t_1^T w_1 & t_1^T w_2 & \dots & t_1^T w_p \\ t_2^T w_1 & \ddots & & \vdots \\ \vdots & & \ddots & \vdots \\ t_n^T w_1 & & & t_n^T w_p \end{bmatrix} \quad (5)$$

Outer Product Representation:

$$TW = \begin{bmatrix} | & | & \dots & | \\ T_1 & T_2 & \dots & T_r \\ | & | & & | \end{bmatrix} \begin{bmatrix} -w_1^T- \\ -w_2^T- \\ \vdots \\ -w_r^T- \end{bmatrix} = \sum_{k=1}^r T_k w_k^T \quad (6)$$

(the sum of rank 1 matrices. TW has rank r if and only if the columns of T are rows of W are linearly independent). In this representation, we can think about T_k as the k th representative taste profile and w_k^T as the k th row of W , or the affinity of each customer with the k th representative profile.

2.2 Linear Independence

Definition 1. (Linear Independence) Vectors $v_1, v_2, \dots, v_n \in \mathbb{R}^p$ are linearly independent vectors if and only if

$$\sum_{j=1}^n \alpha_j v_j = 0 \iff \alpha_j = 0, j = 1, \dots, n \quad (7)$$

Definition 2. (Matrix rank) The rank of a matrix is the maximum number of linearly independent columns. The rank of a matrix is less than the smallest dimension of the matrix.

3 Linear Systems and Vector Norms

Example 2. (Condition on $\text{rank}(A)$ for existence of exact solution)

Consider the linear system of equations $Ax = b$. This means that b is a weighted sum of the columns of A . Suppose A is full rank. Now consider the matrix $[A \ b]$. If the rank of $[A \ b]$ were *greater* than the rank of A (since the number of columns of the matrix increased by 1 and A is assumed full rank, this would imply the rank is $\text{rank}(A) + 1$), this would mean that b could not be written as a linear combination of the columns of A , and that the system would not have an exact solution. Therefore, we must have that $\text{rank}([A \ b]) = \text{rank}(A)$ in order for the system $Ax = b$ to have an exact solution.

To see how the definition of linear independence applies here, observe that $Ax = b \implies Ax - b = 0$. Therefore

$$[A \ b] \begin{bmatrix} x \\ -1 \end{bmatrix} = 0 \quad (8)$$

Thus, if $Ax = b$ has an exact solution, then $[A \ b]$ does not have linearly independent columns.

Example 3. (Condition on $\text{rank}(A)$ for more than one exact solution)

If the system of linear equations $Ax = b$ has more than one exact solution, then there is at least one non zero vector w for which $x + w$ is also a solution. That is, $A(x + w) = b$. If x is an exact solution, then $Ax = b$. This implies $Aw = 0$. Therefore, the columns of A are linearly dependent. Thus, if $\text{rank}(A) < \dim(x)$, then there will be more than one exact solution.

Example 4. (Apply the above conditions) Let

$$A = \begin{bmatrix} 1 & -2 \\ -1 & 2 \\ -2 & 4 \end{bmatrix}, \quad b = \begin{bmatrix} 2 \\ -2 \\ -4 \end{bmatrix} \quad (9)$$

We want to solve $Ax = b$.

- This system has an exact solution, since $\text{rank}(A) = \text{rank}([A \ b])$. This follows since the columns of A are linearly dependent, so it has rank 1, and b is a multiple of the columns of A , so the rank of $[A \ b]$ is also 1.
- Note that $1 = \text{rank}(A) < \dim(x) = 2$. Therefore this system does not have a unique solution.

3.1 Vector Norms

Definition 3. (Vector Norm) A vector norm is a function $\|\cdot\|$ mapping from $\mathbb{R}^n \rightarrow \mathbb{R}$ with the following properties.

1. $\|x\| \geq 0$ for all $x \in \mathbb{R}^n$.

2. $\|x\| = 0$ if and only if $x = 0$.
3. $\|\alpha x\| = |\alpha| \|x\|$ for all $\alpha \in \mathbb{R}, x \in \mathbb{R}^n$.
4. $\|x + y\| \leq \|x\| + \|y\|$ for all $x, y \in \mathbb{R}^n$.

Helpful fact: $\|x\|_{q'} \leq \|x\|_q$ if $1 \leq q \leq q' \leq \infty$.




Figure 1: The l_p norm in \mathbb{R}^2

4 Least Squares

We are given:

1. Vector of labels $y \in \mathbb{R}^n$
2. Matrix of features $X \in \mathbb{R}^{n \times p}$

We want to find:

1. Vector of weights $w \in \mathbb{R}^p$

Assumptions:

1. $n \geq p$, and $\text{rank}(X) = p$.

If $y = Xw$, then we have a system of n linear equations, where the i th equation is

$$y_i = w_1 x_{i1} + w_2 x_{i2} + \cdots + w_p x_{ip} = \sum_{j=1}^p w_j x_{ij} = \langle w, x_{.i} \rangle \quad (10)$$

where $x_{.i}$ is the i th row of X .

In general, $y \neq Xw$ for any w . We define a residual $r_i = y_i - \langle w, x_{.i} \rangle$. Our goal is then to find w $\sum_{i=1}^n |r_i|^2$ (the sum of square residuals/errors).

Why should we minimize the sum of square errors?

1. Magnifies the effect of large errors

2. Allows us to compute derivatives
3. Simple geometric interpretation
4. Coincides with modeling $y = Xw + \epsilon$, where ϵ is Gaussian noise

4.1 Geometric Approach

We know $\hat{r} = y - X\hat{w}$ is orthogonal to the span of the columns of X . Thus $x_i^T \hat{r} = 0$, or $X^T \hat{r} = 0$. This implies $X^T(y - X\hat{w}) = 0$. Thus \hat{w} is a solution to the linear system of equations

$$X^T X \hat{w} = X^T y \quad (11)$$




Figure 2: Geometry of LS in \mathbb{R}^2

Observations:

- The question we're trying to answer: What is the point in X that has the shortest distance to y ? In \mathbb{R}^2 , what are the weights β_1 and β_2 such that $\beta_1 x_1 + \beta_2 x_2$ has the shortest distance to y ?
- $col X$ is the space of all vectors that can be written as $\alpha x_1 + \beta x_2$ for some $\alpha, \beta \in \mathbb{R}$, that is the span of the columns of X . y may not lie in this space.
- The residual vector will form a right angle with $col X$, because any other angle would correspond to a longer distance.

4.2 Vector Calculus Approach

4.2.1 Review of Vector Calculus

Let w be a p -vector and let f be a function of w that maps \mathbb{R}^p to \mathbb{R} . Then the gradient of f with respect to w is

$$\nabla_w f(w) = \begin{pmatrix} \frac{\partial f(w)}{\partial w_1} \\ \vdots \\ \frac{\partial f(w)}{\partial w_p} \end{pmatrix} \quad (12)$$

Example 5. (Gradient of an Inner Product) Let $f(w) = \langle a, w \rangle = w^T a = \sum_{i=1}^n w_i a_i$. Then

$$\nabla_w w^T a = \begin{pmatrix} a_1 \\ a_2 \\ \vdots \\ a_p \end{pmatrix} = a \quad (13)$$

Example 6. (Gradient of an Inner Product, Squared) Let $f(w) = \|w\|^2 = w^T w = w_1^2 + \dots + w_p^2$. Then

$$\nabla_w w^T w = \begin{pmatrix} 2w_1 \\ 2w_2 \\ \vdots \\ 2w_p \end{pmatrix} = 2w \quad (14)$$

(This is a special case of the Quadratic Form discussed below, where $w^T Q w$, and $Q = I$)

Example 7. (Gradient of a Quadratic Form) Let $x \in \mathbb{R}^n$ and $f(x) = x^T Q x$, where Q is symmetric (if Q isn't symmetric we could replace Q with $\frac{1}{2}(Q + Q^T)$). Then

$$\begin{aligned} f(x) &= x^T Q x \\ &= \sum_{i=1}^n \sum_{j=1}^n x_i Q_{ij} x_j \end{aligned}$$

Therefore

$$[\nabla_x f]_k = \frac{df}{dx_k} = \begin{cases} 2Q_{kk}x_k & i = j = k \\ Q_{kj}x_j & i = k, i \neq j \\ Q_{ik}x_i & j = k, j \neq i \end{cases} \quad (15)$$

Therefore

$$\nabla_x f = (Q + Q^T)x \quad (16)$$

If Q is symmetric, then this equals $2Qx$.

4.2.2 Application to Least Squares

Let $f(w) = \|y - Xw\|_2^2$. Then the least squares problem is

$$\hat{w} = \arg \min_w f(w) \quad (17)$$

We can expand $f(w)$ as

$$\begin{aligned} f(w) &= (y - Xw)^T (y - Xw) \\ &= y^T y - y^T Xw - w^T X^T y + w^T X^T Xw \\ &= y^T y - 2w^T X^T y + w^T X^T Xw \end{aligned}$$

Then

$$\nabla_w f(w) = -2X^T y + 2X^T Xw$$

At an optimum we have that \hat{w} solves $X^T y = X^T X\hat{w}$. Then if $(X^T X)^{-1}$ exists, we have that

$$\hat{w} = (X^T X)^{-1} X^T y \quad (18)$$

Theorem 1. (*Sufficient Condition for Existence/Uniqueness of LS Solution*) If the columns of X are linearly independent, then $X^T X$ is non-singular, and there exists a unique least squares solution $\hat{w} = (X^T X)^{-1} X^T y$.

Proof. □

4.3 Positive Definite Matrices

Definition 4 (Positive Definite, pd). A matrix Q ($n \times n$) is positive definite (written $Q \succ 0$) if $x^T Q x > 0$ for all $x \in \mathbb{R}^n$, $x \neq 0$.

Definition 5 (Positive Semi-Definite, psd). A matrix Q ($n \times n$) is positive semi-definite (written $Q \succeq 0$) if $x^T Q x \geq 0$ for all $x \in \mathbb{R}^n$, $x \neq 0$.

Properties of Positive Definite matrices:

1. If $P \succ 0$ and $Q \succ 0$, then $P + Q \succ 0$.
2. If $P \succ 0$ and $\alpha > 0$, then $\alpha P \succ 0$.
3. For any matrix A , $A^T A \succeq 0$ and $AA^T \succeq 0$. Further, if the columns of A are linearly independent, then $A^T A \succ 0$.
4. If $A \succ 0$, then A^{-1} exists.
5. Notation: $A \succ B$ means $A - B \succ 0$.

Example 8. Let

$$X = \begin{pmatrix} 1 & 1 \\ 1 & 1 \\ 1 & 1 \end{pmatrix} \quad (19)$$

Then

$$X^T X = \begin{pmatrix} 3 & 3 \\ 3 & 3 \end{pmatrix} \quad (20)$$

Consider the vector $a = \begin{pmatrix} 1 \\ -1 \end{pmatrix}$. Then $a^T X^T X a = 0$. Therefore $X^T X$ is not positive definite.

4.4 Subspaces

Definition 6. (Subspace) A set of points $S \subseteq \mathbb{R}^n$ is a subspace if

1. $0 \in S$ (S contains the origin)
2. If $x, y \in S$, then $x + y \in S$
3. If $x \in S, \alpha \in \mathbb{R}$, then $\alpha x \in S$.

4.5 Least Squares with Orthonormal Basis for Subspace

Suppose are given a training sample $\{x_i, y_i\}_{i=1}^n, x_i \in \mathbb{R}^p$ and $y \in \mathbb{R}$. If the columns of X (the data matrix) are linearly dependent, then $X^T X$ is not invertible. It is then impossible to tell which features are significant predictors of y .

Given $X = [x_1 \dots x_p]$, the following are options to represent the corresponding subspace spanned by the columns of X :

1. Use X , but then LS can be hard to interpret.
2. Use an orthonormal basis for the subspace.

4.5.1 Orthogonal Matrices and Orthonormal Basis

Definition 7. (Orthonormal basis for X) An orthonormal basis for the columns of X is a collection of vectors $\{u_1, \dots, u_r\}$ such that the span of the columns of X equals the span of $\{u_1, \dots, u_r\}$. That is, $\text{span}(\{x_1, \dots, x_p\}) = \text{span}(\{u_1, \dots, u_r\})$. Furthermore,

$$u_i^T u_j = \begin{cases} 0 & i \neq j \\ 1 & i = j \end{cases} \quad (21)$$

That is, the u vectors are orthogonal and have norm 1.

Observations:

- The rank r of the subspace must satisfy $r \leq \min(n, p)$. r is the number of linearly independent columns of X .
- We can place the basis vectors into a basis matrix $U \in \mathbb{R}^{n \times r}$.

Claim 1. (Properties of orthogonal (basis) matrices) Let $U \in \mathbb{R}^{n \times r}$ be an orthogonal (basis) matrix.

1. $U^T U = I$
2. If U and V are both orthogonal, then UV is also orthogonal.
3. U is length preserving: $\|Uv\|_2 = \|v\|_2$ for $v \in \mathbb{R}^n$.

Proof. We prove each item as follows:

1. We can easily see this from the inner product interpretation of matrix multiplication.
2. $(UV)^T UV = V^T U^T UV = V^T V = I$.
3. $\|Uv\|_2^2 = (Uv)^T Uv = v^T U^T Uv = v^T v = \|v\|_2^2$.

□

4.5.2 Back to LS

Suppose U is an orthonormal basis matrix for our data matrix X . Then, the least-squares problem is

$$\hat{v} = \arg \min_v \|y - Uv\|_2^2 \quad (22)$$

We need \hat{v} to satisfy $U^T U \hat{v} = U^T y$. Thus, $\hat{v} = U^T y$.

4.5.3 Gram-Schmidt Orthogonalization Algorithm

How can we take X and get an orthonormal basis U ?

1. Input $X = [x_1 \dots x_p] \in \mathbb{R}^{n \times p}$
Output: $U = [u_1 \dots u_r] \in \mathbb{R}^{n \times r}$
where $r = \text{rank}(X) \leq \min(n, p)$
2. Initialize $u_1 = \frac{x_1}{\|x_1\|_2}$
3. For $j = 2, 3, \dots, p$
 $x'_j =$ all the components of x_j not represented by u_1, \dots, u_{j-1} .

$$x'_j = x_j - \sum_{i=1}^{j-1} (u_i^T x_j) u_i \quad (23)$$

here $(u_i^T x_j)$ is the least squares weight for u_i .

$$u_j = \begin{cases} \frac{x'_j}{\|x'_j\|_2} & x'_j \neq 0 \\ 0 & x'_j = 0 \end{cases} \quad (24)$$

5 Least Squares Classification

We are given a training sample $\{x_i, y_i\}_{i=1}^n$, $x_i \in \mathbb{R}^p$ and $y \in \mathbb{R}$ (or $y \in \{+1, -1\}$).

Definition 8. (Linear Predictor) We have a linear predictor if each label is a linear combination of the features i.e. we can find weights $\{w_i\}_{i=1}^p$ such that

$$y_i = w_1 x_{i1} + w_2 x_{i2} + \dots + w_p x_{ip} \quad (25)$$

In words, this says the label for observation i is a linear combination of the features for example i .

The steps to complete least squares classification in this environment are as follows:

1. Build a data matrix or feature matrix and label vector

$$X = \begin{bmatrix} -x_1^T - \\ -x_2^T - \\ \vdots \\ -x_n^T - \end{bmatrix} = \begin{bmatrix} x_1^T & 1 \\ x_2^T & 1 \\ \vdots & \\ x_n^T & 1 \end{bmatrix} \in \mathbb{R}^{n \times p}, \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \quad (26)$$

The linear model is then $\hat{y} = Xw$.

2. Solve a least squares optimization problem

$$\hat{w} = \arg \min_w \|y - Xw\|_2^2 = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 \quad (27)$$

(this last equality makes it clear that we are minimizing the sum of squared residuals). If the columns of X are linearly independent, then $X^T X$ is positive definite. Therefore $X^T X$ is invertible. In sum, if $X^T X$ is positive definite, then there exists a unique LS solution

$$\hat{w} = (X^T X)^{-1} X^T y \quad (28)$$

The predicted labels are

$$\begin{aligned} \hat{y} &= Xw \\ &= X(X^T X)^{-1} X^T y \end{aligned}$$

6 Tikhonov Regularization/Ridge Regression

We are given $X \in \mathbb{R}^{n \times p}$ (n training samples, p features) and $y \in \mathbb{R}^n$ (n labels). Our model is $y \approx Xw$, which means $y_i \approx x_i^T w$ for some $w \in \mathbb{R}^p$.

The LS problem is

$$\hat{w}_{LS} = \arg \min_w \|y - Xw\|_2^2 = \arg \min_w \sum_{i=1}^n (y_i - x_i^T w)^2 \quad (29)$$

There are two cases

1. If X is full rank (i.e. the columns of X are linearly independent), then \hat{w}_{LS} is unique and

$$\hat{w}_{LS} = (X^T X)^{-1} X^T y \quad (30)$$

2. If X is not full rank, then $X^T X$ is not invertible. \hat{w}_{LS} is not unique; there are infinitely many solutions.

6.1 Tikhonov Regularization Derivation

In this second case, we can define a new objective

$$\hat{w} = \arg \min_w \|y - Xw\|_2^2 + \lambda \|w\|_2^2 \quad (31)$$

where $\|y - Xw\|_2^2$ measures the fit to the data, $\lambda > 0$ is a regularization parameter or tuning parameter, and $\|w\|_2^2$ is a regularizer. $\|w\|_2^2$ measures the energy in w .

Observations about this problem:

1. \hat{w} is unique even when no unique least square solution exists
2. Even when X is full rank, $X^T X$ can be badly behaved, and regularization adjusts for this.

6.1.1 Derivation with Vector Calculus

Let $f(w) = \|y - Xw\|_2^2 + \lambda \|w\|_2^2$. Then

$$\begin{aligned} f(w) &= y^T y - 2w^T X^T y + w^T X^T X w + \lambda w^T w \\ &= y^T y - 2w^T X^T y + w^T (X^T X + \lambda I) w \end{aligned}$$

Then

$$\nabla_w f(w) = -2X^T y + 2(X^T X + \lambda I) \quad (32)$$

If $(X^T X + \lambda I)$ is invertible, then $\hat{w} = (X^T X + \lambda I)^{-1} X^T y$. BUT, $(X^T X + \lambda I)$ is *always* invertible. Recall that if a matrix is positive definite, then it is invertible. We can show that

$(X^T X + \lambda I)$ is indeed positive definite and hence invertible. To see this, fix $0 \neq a \in \mathbb{R}^n$, then

$$\begin{aligned} a^T (X^T X + \lambda I) a &= a^T X^T X a + \lambda a^T a \\ &= \|Xa\|_2^2 + \lambda \|a\|_2^2 \end{aligned}$$

Now note that $\|Xa\|_2^2 \geq 0$ (it could be 0 if X is not full rank and a is in the null space of X – this is what causes troubles with LS) but $\lambda \|a\|_2^2 > 0$. Therefore, $(X^T X + \lambda I)$ is positive definite.

6.1.2 Alternative Derivation

Note that for vectors a, b ,

$$\|a\|_2^2 + \|b\|_2^2 = \left\| \begin{bmatrix} a \\ b \end{bmatrix} \right\|_2^2 \quad (33)$$

Therefore,

$$\begin{aligned} f(w) &= \|y - Xw\|_2^2 + \lambda \|w\|_2^2 \\ &= \|y - Xw\|_2^2 + \left\| \sqrt{\lambda} w \right\|_2^2 \\ &= \left\| \begin{bmatrix} y - Xw \\ \sqrt{\lambda} w \end{bmatrix} \right\|_2^2 \\ &= \left\| \begin{bmatrix} y \\ 0 \end{bmatrix} - \begin{bmatrix} Xw \\ \sqrt{\lambda} w \end{bmatrix} \right\|_2^2 \\ &= \left\| \begin{bmatrix} y \\ 0 \end{bmatrix} - \begin{bmatrix} X \\ \sqrt{\lambda} I \end{bmatrix} w \right\|_2^2 \\ &= \|\tilde{y} - \tilde{X}w\|_2^2 \end{aligned}$$

We can solve this problem with LS, so that

$$\hat{w} = (\tilde{X}^T \tilde{X})^{-1} \tilde{X}^T \tilde{y} \quad (34)$$

where

$$\tilde{X}^T \tilde{X} = X^T X + \lambda I \quad (35)$$

and

$$\tilde{X} \tilde{y} = X^T y \quad (36)$$

Thus this is equivalent to the derivation above.

7 Singular Value Decomposition